

Cover Letter Cum Research Background

July 8, 2026

Dear Prof,

I am Adarsh, currently a CBRAIN Research Intern at the Centre for Brain Research (CBR), Indian Institute of Science (IISc), and an M.Sc. graduate in Data Science with a background in Applied Mathematics and Statistics (B.Sc. in Math & Stats). Currently at CBR, I work on machine learning methods for dementia research using multimodal longitudinal clinical data, with recent work spanning longitudinal risk prediction, survival analysis, and multimodal representation learning. Previously, I was an ML Research Intern at the Indian Institute of Information Technology (IIIT) Kottayam, where I developed vision-language models for plant disease diagnosis by integrating CLIP-based visual representations with large language models, gaining experience in building and evaluating multimodal systems.

My master’s research investigates how neural representations align with latent spaces of modern generative models, with a focus on identifying where linear mappings fail and what this reveals about representation mismatch in brain-to-image reconstruction. This work lies at the intersection of NeuroAI and generative modeling, where I aim to understand what the limitations of these models reveal about representation learning. I have also developed strong hands-on proficiency in the deep learning stack. As outlined in the project description and to align with my pertinent experience, the following project summaries illustrate the diversity of my completed and ongoing research projects.

My MSc thesis examines how representations align across neural systems and generative models through two complementary studies.

(1) Brain2VLM: Hierarchical Alignment Between Cortical Representations and Vision-Language Latent Spaces (1) This work investigates how neural activity aligns with latent spaces of diffusion-based vision-language models for brain-to-image reconstruction. While prior work demonstrates strong reconstruction performance using latent diffusion models (2; 3), the structure of the brain-to-latent mapping remains poorly understood. I evaluate whether linear mappings are sufficient for decoding different types of representations, using ridge regression and a nonlinear residual MLP on fMRI data from the Natural Scenes Dataset (4). Prior work suggests hierarchical organization of visual and semantic representations (5), motivating this analysis. I show that linear mappings are sufficient for structural diffusion latents but fail for semantic embeddings, where nonlinear decoding yields substantial improvements ($\Delta corr \approx 0.47$ vs. 0.05–0.06). However, increased decoder expressivity introduces distribution shifts that degrade generative compatibility, revealing a trade-off between alignment accuracy and reconstruction fidelity. Despite its simplicity, this approach achieves competitive reconstruction performance ($PixCorr \approx 0.33$, $CLIP \approx 85\%$) relative to recent diffusion-based approaches such as Brain-Diffuser and MindLDM (3; 6). This suggests that the brain-to-latent interface is a primary bottleneck in current decoding systems, highlighting a fundamental mismatch between neural representations and generative model latent spaces. A glimpse of the results (Brain2VLM) is shown in Figure 1 below.

(2) SkinGenBench: Generative Model and Preprocessing Effects for Synthetic Dermoscopic Augmentation in Melanoma Diagnosis. (7). This work addresses the lack of reproducible, clinically grounded benchmarks for




Figure 1: in the row, cortical surface based brain maps of alignment differences; PCA of ground-truth (GT), linearly decoded and MLP decoded CLIP latents; GT (top), linear decoder (middle), and MLP (bottom) based samples of good, average, and challenging reconstructions.

synthetic dermoscopic data (8; 9), particularly in underrepresented melanoma cases, which remain scarce in public datasets such as HAM10000 (10) and ISIC (8). I evaluate how generative model choice (GANs vs. diffusion models) and preprocessing complexity affect both image quality and downstream diagnostic performance, comparing StyleGAN2-ADA and diffusion models (11). I find that StyleGAN2-ADA consistently outperforms diffusion models in fidelity ($FID \approx 65.5$, $KID \approx 0.05$) and diagnostic utility, improving melanoma F1 by 8–15%, while diffusion models exhibit higher variance but reduced perceptual coherence, consistent with recent comparative findings. (12) This challenges the assumption that diffusion models are universally superior, showing that in low-data clinical settings, GAN-based approaches remain more reliable for diagnostic tasks. A glimpse of the results (SkinGenBench) is shown in Figure 2 below.



Figure 2: In the row, t-SNE embeddings of synthesized mel samples, FID, IS, and KID scores (epochwise, throughout training), Grad-CAM visualisation (Top 2 best performing classifiers), Sample input and output of the artifact removal algorithm (advanced preprocessing)

My work is driven by the following research questions: (1) How do neural representations align with latent spaces of generative models, and where do linear mappings fail? (2) How does decoder expressivity influence alignment and generative compatibility in brain-to-image reconstruction? (3) How can representation learning and multimodal AI improve clinically meaningful prediction and decision-making in healthcare?

I have delivered invited seminars on ML research methodology, covering experimental design, reproducibility, and practical workflows. I have implemented core deep learning components from first principles, including a transformer-based language model (13), with technical blogs on tokenization and attention (14; 15). My applied work includes an end-to-end life expectancy modeling pipeline using WHO data, including multicollinearity analysis (16), and a sign language translation system integrating computer vision and NLP, presented at an IEEE conference (17).

I see this opportunity as a natural continuation of my research interests in representation learning, deep learning, and multimodal AI. My background in implementing and extending modern deep learning architectures, computer vision, multimodal learning, and explainable AI has prepared me well to contribute to these research directions. Through my research projects, I have gained hands-on experience with PyTorch, deep learning experimentation, model implementation from research papers, reproducible machine learning pipelines, and training models on cloud GPU infrastructure. I am excited by the opportunity to contribute to your lab while further developing my research skills under your supervision.

Thanks for considering my application!

References

- [1] N. A. Adarsh Pritam, J. S. O, and S. Jain, “**Brain2VLM: Hierarchical Alignment Between Cortical Representations and Vision-Language Latent Spaces**,” 2026. [Online]. Available: <https://doi.org/10.13140/RG.2.2.26951.05288>
- [2] Y. Takagi and S. Nishimoto, “High-resolution image reconstruction with latent diffusion models from human brain activity,” in *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 14 453–14 463.
- [3] M. Ferrante, T. Boccato, F. Ozcelik, R. VanRullen, and N. Toschi, “Generative Multimodal Decoding: Reconstructing Images and Text from Human fMRI,” in *Deep Generative Models for Health Workshop NeurIPS*, 2023.
- [4] E. J. Allen, G. St-Yves, Y. Wu, J. L. Breedlove, J. S. Prince, L. T. Dowdle, M. Nau, B. Caron, F. Pestilli, I. Charest, J. B. Hutchinson, T. Naselaris, and K. Kay, “A massive 7T fMRI dataset to bridge cognitive neuroscience and artificial intelligence,” *Nature Neuroscience*, vol. 25, no. 1, pp. 116–126, 2022.
- [5] A. G. Huth, S. Nishimoto, A. T. Vu, and J. L. Gallant, “A Continuous Semantic Space Describes the Representation of Thousands of Object and Action Categories across the Human Brain,” *Neuron*, vol. 76, no. 6, pp. 1210–1224, 2012.
- [6] J. Guo, C. Yi, F. Li, P. Xu, and Y. Tian, “MindLDM: Reconstruct Visual Stimuli from fMRI Using Latent Diffusion Model,” in *2024 IEEE International Conference on Computational Intelligence and Virtual Environments for Measurement Systems and Applications (CIVEMSA)*, 2024, pp. 1–6.
- [7] N. A. Adarsh Pritam, J. S. O, and S. Jain, “**SkinGenBench: Generative Model and Preprocessing Effects for Synthetic Dermoscopic Augmentation in Melanoma Diagnosis**,” <https://arxiv.org/abs/2512.17585>, 2025.
- [8] B. Cassidy, C. Kendrick, A. Brodzicki, J. Jaworek-Korjakowska, and M. H. Yap, “Analysis of the ISIC image datasets: Usage, benchmarks and recommendations,” *Medical Image Analysis*, vol. 75, p. 102305, 2022.
- [9] T. Wang, Y. Lei, Y. Fu, J. F. Wynne, W. J. Curran, T. Liu, and X. Yang, “A review on medical imaging synthesis using deep learning and its clinical applications,” *J Appl Clin Med Phys*, vol. 22, no. 1, pp. 11–36, 2021.
- [10] P. Tschandl, C. Rosendahl, and H. Kittler, “The HAM10000 dataset, a large collection of multi-source dermatoscopic images of common pigmented skin lesions,” *Sci Data*, vol. 5, no. 1, p. 180161, 2018.
- [11] J. Ho, A. Jain, and P. Abbeel, “Denoising Diffusion Probabilistic Models,” vol. 33, pp. 6840–6851, 2020.
- [12] W. Wang, J. Xia, G. Luo, S. Dong, X. Li, J. Wen, and S. Li, “Diffusion model for medical image denoising, reconstruction and translation,” *Computerized Medical Imaging and Graphics*, vol. 124, p. 102593, 2025.
- [13] N. A. Adarsh Pritam, “**LLaMA LLM from Scratch**,” <https://github.com/adarsh-crafts/llama-llm-from-scratch>, 2024.
- [14] —, “**Build a Byte Pair Encoding (BPE) Tokenizer from Scratch in Python**,” <https://medium.com/@adarsh-ai/build-a-byte-pair-encoding-bpe-tokenizer-from-scratch-in-python-0dc32c6410f7>, 2024.
- [15] —, “**A Beginner’s Guide to Multi-Head Self-Attention in LLMs**,” <https://medium.com/@adarsh-ai/a-beginners-guide-to-multi-head-self-attention-in-llms-1a4ea8be6fb2>, 2024.
- [16] —, “**WHO Life Expectancy Modeling: Socioeconomic and Health Predictors using Interpretable Machine Learning**,” <https://github.com/adarsh-crafts/life-expectancy-modeling>, 2026.
- [17] Adarsh Pritam, N. A. and A. Kurian, “**An End-to-End Sign Language Translation Pipeline from Static Gestures to English Using T5**,” in *2025 International Conference on Emerging Technologies in Computing and Communication (ETCC)*, 2025, pp. 1–6.